



CSE5014 CRYPTOGRAPHY AND NETWORK SECURITY

Dr. QI WANG

Department of Computer Science and Engineering

Office: Room413, CoE South Tower

Email: wangqi@sustech.edu.cn

Perfect secrecy

- **Definition 1.5** An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if for every distribution over \mathcal{M} , every $m \in \mathcal{M}$, and every $c \in \mathcal{C}$ with $\Pr[C = c] > 0$, it holds that
$$\Pr[M = m | C = c] = \Pr[M = m]$$



Perfect secrecy

- **Definition 1.5** An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if **for every distribution over \mathcal{M} , every $m \in \mathcal{M}$, and every $c \in \mathcal{C}$ with $\Pr[C = c] > 0$, it holds that**

$$\Pr[M = m | C = c] = \Pr[M = m]$$

Equivalently, for every set $M \subseteq \{0, 1\}^\ell$ of plaintexts, and for every strategy used by Eve, if we choose at random $x \in M$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x after seeing $Enc_k(x)$ is **at most $1/|M|$** , i.e.,

$$\Pr[Eve(Enc_k(x)) = x] \leq 1/|M|$$



Perfect secrecy

- Another two **equivalent** definitions



Perfect secrecy

- Another two **equivalent** definitions

Definition 1.6 *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if **for every two distinct plaintexts** $\{x_0, x_1\} \in \mathcal{M}$, and for every strategy used by Eve, **if we choose at random** $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x_b after seeing the ciphertext $c = Enc_k(x_b)$ is **at most** $1/2$.



Perfect secrecy

- Another two **equivalent** definitions

Definition 1.6 *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if **for every two distinct plaintexts** $\{x_0, x_1\} \in \mathcal{M}$, and for every strategy used by Eve, **if we choose at random** $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x_b after seeing the ciphertext $c = Enc_k(x_b)$ is **at most** $1/2$.

Definition 1.7 *Perfect secrecy*. Two probability distributions X, Y over $\{0, 1\}^\ell$ are *identical*, denoted by $X \equiv Y$, if for every $y \in \{0, 1\}^\ell$, $\Pr[X = y] = \Pr[Y = y]$. An encryption scheme (Gen, Enc, Dec) is *perfectly secure* if **for every pair of plaintexts** x, x' , we have $Enc_{U_n}(x) \equiv Enc_{U_n}(x')$.

Perfect secrecy

- Another two **equivalent** definitions

Definition 1.6 *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if **for every two distinct plaintexts** $\{x_0, x_1\} \in \mathcal{M}$, and for every strategy used by Eve, **if we choose at random** $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x_b after seeing the ciphertext $c = Enc_k(x_b)$ is **at most** $1/2$.

Definition 1.7 *Perfect secrecy*. Two probability distributions X, Y over $\{0, 1\}^\ell$ are *identical*, denoted by $X \equiv Y$, if for every $y \in \{0, 1\}^\ell$, $\Pr[X = y] = \Pr[Y = y]$. An encryption scheme (Gen, Enc, Dec) is *perfectly secure* if **for every pair of plaintexts** x, x' , we have $Enc_{U_n}(x) \equiv Enc_{U_n}(x')$.

Q: Does this mean that for **every** k , $Enc_k(x) = Enc_k(x')$?



Perfect security

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{x_0, x_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x_b after seeing the ciphertext $c = Enc_k(x_b)$ is at most $1/2$.

Key point: The ciphertext c reveals zero additional information about the plaintext m .

$$\Pr[M = m | C = c] = \Pr[M = m]$$

Perfect security

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{x_0, x_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x_b after seeing the ciphertext $c = Enc_k(x_b)$ is at most $1/2$.

Key point: The ciphertext c reveals **zero additional information** about the plaintext m .

$$\Pr[M = m | C = c] = \Pr[M = m]$$

Equivalently, for every set $M \subseteq \{0, 1\}^\ell$ of plaintexts, and for every strategy used by Eve, if we choose at random $x \in M$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x after seeing $Enc_k(x)$ is **at most $1/|M|$** , i.e.,

$$\Pr[Eve(Enc_k(x)) = x] \leq 1/|M|$$



Perfect Secrecy

- **Theorem 1.8** ([Two-to-Many Theorem](#)) The scheme (Gen, Enc, Dec) is *perfectly secure* if and only if $Pr[Eve(Enc_k(x_b)) = x_b] \leq 1/2$.



Perfect Secrecy

- **Theorem 1.8** ([Two-to-Many Theorem](#)) The scheme (Gen, Enc, Dec) is *perfectly secure* if and only if
$$Pr[Eve(Enc_k(x_b)) = x_b] \leq 1/2.$$

Proof.

Perfect Secrecy

- **Theorem 1.8** ([Two-to-Many Theorem](#)) The scheme (Gen, Enc, Dec) is *perfectly secure* if and only if
$$Pr[Eve(Enc_k(x_b)) = x_b] \leq 1/2.$$

Proof.

The “only if” part is easy (by definition, this is the [special case](#) that $|M| = 2$).

The “if” part is tricky.



Perfect Secrecy

- **Theorem 1.8 (Two-to-Many Theorem)** The scheme (Gen, Enc, Dec) is *perfectly secure* if and only if
$$Pr[Eve(Enc_k(x_b)) = x_b] \leq 1/2.$$

Proof.

The “only if” part is easy (by definition, this is the *special case* that $|M| = 2$).

The “if” part is tricky.

We need to show that if there is some set M and some strategy for Eve to guess a plaintext chosen from M with probability larger than $1/|M|$, then there is also some set M' of size 2 and a strategy Eve' for Eve to guess a plaintext chosen from M' with probability larger than $1/2$.

We fix $x_0 = 0^\ell$ and pick x_1 at random in M . Then it holds that for random key k and message $x_1 \in M$,

$$Pr_{k \leftarrow \{0,1\}^n, x_1 \leftarrow M}[Eve(Enc_k(x_1)) = x_1] > 1/|M|.$$

On the other hand, for every choice of k , $x' = Eve(Enc_k(x_0))$ is a fixed string independent on the choice of x_1 , and so if we pick x_1 at random in M , then the probability that $x_1 = x'$ is at most $1/|M|$, i.e.,

$$Pr_{k \leftarrow \{0,1\}^n, x_1 \leftarrow M}[Eve(Enc_k(x_0)) = x_1] \leq 1/|M|.$$

Due to the linearity of expectation, there exists some x_1 satisfying

$$Pr[Eve(Enc_k(x_1)) = x_1] > Pr[Eve(Enc_k(x_0)) = x_1]. \text{ (why?)}$$

We now define a new attacker Eve' as: $Eve'(c) = \begin{cases} x_1, & \text{if } Eve(c) = x_1 \\ x_i, i \in \{0, 1\} \text{ at random,} & \text{otherwise} \end{cases}$

This means the probability that $Eve'(Enc_k(x_b)) = x_b$ is larger than $1/2$ (Why?).



One-time Pad

- The *XOR* operation: $a \oplus b = a + b \bmod 2$.



One-time Pad

- The *XOR* operation: $a \oplus b = a + b \bmod 2$.

$$a \oplus 0 = a$$

$$a \oplus a = 0$$

$$a \oplus b = b \oplus a \text{ (Commutativity)}$$

$$a \oplus (b \oplus c) = (a \oplus b) \oplus c \text{ (Associativity)}$$



One-time Pad

- The *XOR* operation: $a \oplus b = a + b \bmod 2$.

$$a \oplus 0 = a$$

$$a \oplus a = 0$$

$$a \oplus b = b \oplus a \text{ (Commutativity)}$$

$$a \oplus (b \oplus c) = (a \oplus b) \oplus c \text{ (Associativity)}$$

The One-time Pad scheme (Vernam 1917, Shannon 1949):

$$n = |k| = |x|, \text{ Enc} : \{0, 1\}^n \times \{0, 1\}^n \rightarrow \{0, 1\}^n$$

$$\text{Enc}_k(x) = x \oplus k$$

$$\text{Dec}_k(y) = y \oplus k$$



One-time Pad

- The *XOR* operation: $a \oplus b = a + b \bmod 2$.

$$a \oplus 0 = a$$

$$a \oplus a = 0$$

$$a \oplus b = b \oplus a \text{ (Commutativity)}$$

$$a \oplus (b \oplus c) = (a \oplus b) \oplus c \text{ (Associativity)}$$

The One-time Pad scheme (Vernam 1917, Shannon 1949):

$$n = |k| = |x|, \text{ Enc} : \{0, 1\}^n \times \{0, 1\}^n \rightarrow \{0, 1\}^n$$

$$\text{Enc}_k(x) = x \oplus k$$

$$\text{Dec}_k(y) = y \oplus k$$

Validity:

$$\text{Dec}_k(\text{Enc}_k(x)) = (x \oplus k) \oplus k = x \oplus (k \oplus k) = x \oplus 0^n = x$$



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Proof. Fix **arbitrary** distribution over $\mathcal{M} = \{0, 1\}^n$, and arbitrary $m, c \in \{0, 1\}^n$.



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Proof. Fix **arbitrary** distribution over $\mathcal{M} = \{0, 1\}^n$, and arbitrary $m, c \in \{0, 1\}^n$.

$$\Pr[M = m \mid C = c] = ?$$



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Proof. Fix **arbitrary** distribution over $\mathcal{M} = \{0, 1\}^n$, and arbitrary $m, c \in \{0, 1\}^n$.

$$\begin{aligned}\Pr[M = m \mid C = c] &=? \\ &= \Pr[C = c \mid M = m] \cdot \Pr[M = m] / \Pr[C = c]\end{aligned}$$



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Proof. Fix **arbitrary** distribution over $\mathcal{M} = \{0, 1\}^n$, and arbitrary $m, c \in \{0, 1\}^n$.

$$\begin{aligned}\Pr[M = m \mid C = c] &=? \\ &= \Pr[C = c \mid M = m] \cdot \Pr[M = m] / \Pr[C = c]\end{aligned}$$

$$\begin{aligned}\Pr[C = c] &= \sum_{m'} \Pr[C = c \mid M = m'] \cdot \Pr[M = m'] \\ &= \sum_{m'} \Pr[K = m' \oplus c] \cdot \Pr[M = m'] \\ &= \sum_{m'} 2^{-n} \cdot \Pr[M = m'] \\ &= 2^{-n}\end{aligned}$$

Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Proof. Fix **arbitrary** distribution over $\mathcal{M} = \{0, 1\}^n$, and arbitrary $m, c \in \{0, 1\}^n$.

$$\begin{aligned}\Pr[M = m \mid C = c] &=? \\&= \Pr[C = c \mid M = m] \cdot \Pr[M = m] / \Pr[C = c] \\&= \Pr[K = m \oplus c] \cdot \Pr[M = m] / 2^{-n} \\&= 2^{-n} \cdot \Pr[M = m] / 2^{-n} \\&= \Pr[M = m]\end{aligned}$$

$$\begin{aligned}\Pr[C = c] &= \sum_{m'} \Pr[C = c \mid M = m'] \cdot \Pr[M = m'] \\&= \sum_{m'} \Pr[K = m' \oplus c] \cdot \Pr[M = m'] \\&= \sum_{m'} 2^{-n} \cdot \Pr[M = m'] \\&= 2^{-n}\end{aligned}$$



One-time Pad

- *Q*: Is this the end of cryptography?



One-time Pad

- \mathcal{Q} : Is this the end of cryptography?

We need more:

- ◇ Use the same key for many plaintexts
- ◇ Use n -bit key for $2n$ -bit plaintexts.

One-time Pad

- *Q*: Is this the end of cryptography?

We need more:

- ◇ Use the same key for many plaintexts
- ◇ Use *n -bit key* for *$2n$ -bit plaintexts*.

Theorem 1.10 (*Limitations of perfect secrecy*) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with *n -bit* plaintexts and *$(n - 1)$ -bit* keys.



One-time Pad

- *Q*: Is this the end of cryptography?

We need more:

- ◇ Use the same key for many plaintexts
- ◇ Use *n -bit key* for *$2n$ -bit plaintexts*.

Theorem 1.10 (*Limitations of perfect secrecy*) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with *n -bit* plaintexts and *$(n - 1)$ -bit* keys.

Proof. See blackboard.

One-time Pad

- **Q:** Is this the end of cryptography?

We need more:

- ◇ Use the same key for many plaintexts
- ◇ Use n -bit key for $2n$ -bit plaintexts.

Theorem 1.10 (Limitations of perfect secrecy) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with n -bit plaintexts and $(n - 1)$ -bit keys.

Proof. See blackboard.

Suppose that (Gen, Enc, Dec) is such an encryption scheme. Denote by Y_0 the distribution $Enc_{U_{n-1}}(0^n)$ and by S_0 its support. Since there are only 2^{n-1} possible keys, we have $|S_0| \leq 2^{n-1}$.

Now for every key k the function $Enc_k(\cdot)$ is one-to-one and hence its image is of size $\geq 2^n$. This means that for every k , there exists x such that $Enc_k(x) \notin S_0$. Fix such a k and x , then the distribution $Enc_{U_{n-1}}(x)$ does not have the same support as Y_0 and hence it is not identical to Y_0 .

Key generation

- When describing algorithms, we assume access to **uniformly distributed bits/bytes**. Where do these actually come from?



Key generation

- When describing algorithms, we assume access to **uniformly distributed bits/bytes**. Where do these actually come from?
- Precise details depend on the system
 - Linux or unix: `/dev/random` or `/dev/urandom`
 - **Do not** use `rand()` or `java.util.Random`
 - Use crypto libraries instead

Random-number generation

- Two steps:

1. Continually collect a “pool” of high-entropy (“unpredictable”) data

2. (Smoothing) When random bits are requested, process this data to generate a sequence of uniform, independent bits/bytes



Random-number generation

- Two steps:

1. Continually collect a “pool” of **high-entropy** (“unpredictable”) data
 - Must ultimately come from some physical process (keystroke/mouse movements, delays between network events, hard-disk access times, hardware random-number generation (Intel), etc.)
2. (Smoothing) When random bits are requested, process this data to generate a sequence of uniform, independent bits/bytes



Random-number generation

■ Two steps:

1. Continually collect a “pool” of **high-entropy** (“unpredictable”) data
 - Must ultimately come from some physical process (keystroke/mouse movements, delays between network events, hard-disk access times, hardware random-number generation (Intel), etc.)
2. (Smoothing) When random bits are requested, process this data to generate a sequence of uniform, independent bits/bytes
 - Need to eliminate both *bias* and *dependencies*

Random-number generation

■ Step 2: Smoothing

von Neumann technique for eliminating *bias*:

- Collect two bits per output bit
 - 01 \rightarrow 0, 10 \rightarrow 1, 00, 11 \rightarrow skip
- Note that this assumes *independence* (as well as constant bias)

Random-number generation

- Step 2: Smoothing
 - von Neumann technique for eliminating *bias*:
 - Collect two bits per output bit
 - 01 \rightarrow 0, 10 \rightarrow 1, 00, 11 \rightarrow skip
 - Note that this assumes *independence* (as well as constant bias)
- Read desired number of bytes from “/dev/urandom”

Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Theorem 1.10 (*Limitations of perfect secrecy*) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with n -bit plaintexts and $(n - 1)$ -bit keys.

- The key is **as long as** the message
- **Only secure** if each key is used to encrypt a **single** message
- Trivially **broken by a known-plaintext attack**

Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Theorem 1.10 (*Limitations of perfect secrecy*) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with n -bit plaintexts and $(n - 1)$ -bit keys.

- The key is **as long as** the message
- **Only secure** if each key is used to encrypt a **single** message
- Trivially **broken by a known-plaintext attack**

Q: What about if using the same key twice?



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Theorem 1.10 (*Limitations of perfect secrecy*) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with n -bit plaintexts and $(n - 1)$ -bit keys.

- The key is **as long as** the message
- **Only secure** if each key is used to encrypt a **single** message
- Trivially **broken by a known-plaintext attack**

Q: What about if using the same key twice?

Say $c_1 = k \oplus m_1$ and $c_2 = k \oplus m_2$.

$$c_1 \oplus c_2 = (k \oplus m_1) \oplus (k \oplus m_2) = m_1 \oplus m_2$$



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Theorem 1.10 (*Limitations of perfect secrecy*) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with n -bit plaintexts and $(n - 1)$ -bit keys.

- The key is **as long as** the message
- **Only secure** if each key is used to encrypt a **single** message
- Trivially **broken by a known-plaintext attack**

Q: What about if using the same key twice?

Say $c_1 = k \oplus m_1$ and $c_2 = k \oplus m_2$.

$$c_1 \oplus c_2 = (k \oplus m_1) \oplus (k \oplus m_2) = m_1 \oplus m_2$$

$m_1 \oplus m_2$ **leaks** information about m_1, m_2



Perfect security

- **Theorem 1.9** One-time pad is *perfectly secure*.

Theorem 1.10 (*Limitations of perfect secrecy*) There is **no** *perfectly secure* encryption schemes (Gen, Enc, Dec) with n -bit plaintexts and $(n - 1)$ -bit keys.

- The key is **as long as** the message
- **Only secure** if each key is used to encrypt a **single** message
- Trivially **broken by a known-plaintext attack**

Whenever faced with an impossibility result, it is **a good idea** to examine whether we can **relax** these assumptions to still get what we want (or at least something close to that).



Statistical Security

- Eve has a **tiny** advantage in its posteriori guessing probability compared to the priori probability.



Statistical Security

- Eve has a **tiny** advantage in its posteriori guessing probability compared to the priori probability.

We may say that an encryption scheme is *ϵ -statistically indistinguishable* if the probability that Eve guesses which of the two messages was encrypted is at most $1/2 + \epsilon$.



Statistical Security

- Eve has a **tiny** advantage in its posteriori guessing probability compared to the priori probability.

We may say that an encryption scheme is *ϵ -statistically indistinguishable* if the probability that Eve guesses which of the two messages was encrypted is at most $1/2 + \epsilon$.

Q:

- 1) Is the **relaxed** definition still strong enough in practice?
- 2) Does the relaxation buy something we could not get with the original definition (perfect security)?



Statistical Security

- Eve has a **tiny** advantage in its posteriori guessing probability compared to the priori probability.

We may say that an encryption scheme is *ϵ -statistically indistinguishable* if the probability that Eve guesses which of the two messages was encrypted is at most $1/2 + \epsilon$.

Q:

- 1) Is the **relaxed** definition still strong enough in practice?
- 2) Does the relaxation buy something we could not get with the original definition (perfect security)?

A:

- 1) Yes, if ϵ is small (say 10^{-6} or even 10^{-100})
- 2) No, we **cannot** have key shorter than the message.



ϵ -Statistical Security

- **Definition 2.1** Let X and Y be two distributions over $\{0, 1\}^n$. The *statistical distance* of X and Y , denoted by $\Delta(X, Y)$ is defined to be
$$\max_{T \subseteq \{0, 1\}^n} |\Pr[X \in T] - \Pr[Y \in T]|.$$
If $\Delta(X, Y) \leq \epsilon$, we say that $X \equiv_{\epsilon} Y$.



ϵ -Statistical Security

- **Definition 2.1** Let X and Y be two distributions over $\{0, 1\}^n$. The *statistical distance* of X and Y , denoted by $\Delta(X, Y)$ is defined to be

$$\max_{T \subseteq \{0, 1\}^n} |\Pr[X \in T] - \Pr[Y \in T]|.$$

If $\Delta(X, Y) \leq \epsilon$, we say that $X \equiv_{\epsilon} Y$.

Definition 2.2 *ϵ -Statistical Security*. An encryption scheme (Gen, Enc, Dec) is *ϵ -statistically secure* if for every pair of plaintexts m, m' , we have $Enc_{U_n}(m) \equiv_{\epsilon} Enc_{U_n}(m')$.



■ Lemma 2.3

$$\Delta(X, Y) = \frac{1}{2} \sum_{w \in \text{Supp}(X) \cup \text{Supp}(Y)} |Pr[X = w] - Pr[Y = w]|$$



■ Lemma 2.3

$$\Delta(X, Y) = \frac{1}{2} \sum_{w \in \text{Supp}(X) \cup \text{Supp}(Y)} |Pr[X = w] - Pr[Y = w]|$$

Proof. See blackboard.



■ Lemma 2.3

$$\Delta(X, Y) = \frac{1}{2} \sum_{w \in \text{Supp}(X) \cup \text{Supp}(Y)} |Pr[X = w] - Pr[Y = w]|$$

Proof. See blackboard.

Observations:

$$0 \leq \Delta(X, Y) \leq 1$$

$$\Delta(X, Y) = 0 \text{ if } X = Y$$

$$0 \leq \Delta(X, Y) \leq \Delta(X, Z) + \Delta(Z, Y)$$



■ Lemma 2.3

$$\Delta(X, Y) = \frac{1}{2} \sum_{w \in \text{Supp}(X) \cup \text{Supp}(Y)} |Pr[X = w] - Pr[Y = w]|$$

Proof. See blackboard.

Observations:

$$0 \leq \Delta(X, Y) \leq 1$$

$$\Delta(X, Y) = 0 \text{ if } X = Y$$

$$0 \leq \Delta(X, Y) \leq \Delta(X, Z) + \Delta(Z, Y)$$

Δ is a *metric*.



ϵ -Statistical Security

- **Lemma 2.4** Eve has at most $1/2 + \epsilon$ success probability if and only if for every pair of m_1, m_2 ,
$$\Delta(\text{Enc}_{U_n}(m_1), \text{Enc}_{U_n}(m_2)) \leq 2\epsilon.$$



ϵ -Statistical Security

- **Lemma 2.4** Eve has at most $1/2 + \epsilon$ success probability if and only if for every pair of m_1, m_2 ,
$$\Delta(\text{Enc}_{U_n}(m_1), \text{Enc}_{U_n}(m_2)) \leq 2\epsilon.$$

Proof.

Suppose that Eve has $1/2 + \epsilon$ success probability with m_1, m_2 . Let $p_{i,j} = \Pr[\text{Eve}(\text{Enc}_{U_n}(m_i)) = j]$. Then we have

$$p_{1,1} + p_{1,2} = 1$$

$$p_{2,1} + p_{2,2} = 1$$

$$(1/2)p_{1,1} + (1/2)p_{2,2} \leq 1/2 + \epsilon.$$

The last two together imply that

$$p_{1,1} - p_{2,1} \leq 2\epsilon,$$

which means that if we let T be the set $\{c : \text{Eve}(c) = 1\}$, then T demonstrates that $\Delta(\text{Enc}_{U_n}(m_1), \text{Enc}_{U_n}(m_2)) \leq 2\epsilon$.

Similarly, if we have such a set T , we can define an attacker from it that succeeds with probability $1/2 + \epsilon$.



Limitation of ϵ -Statistical Security

- **Theorem 2.5** Let (Gen, Enc, Dec) be a valid encryption with $Enc : \{0, 1\}^n \times \{0, 1\}^{n+1} \rightarrow \{0, 1\}^*$. Then there exist plaintexts m_1, m_2 with $\Delta(Enc_{U_n}(m_1), Enc_{U_n}(m_2)) > 1/2$.



Limitation of ϵ -Statistical Security

- **Theorem 2.5** Let (Gen, Enc, Dec) be a valid encryption with $Enc : \{0, 1\}^n \times \{0, 1\}^{n+1} \rightarrow \{0, 1\}^*$. Then there exist plaintexts m_1, m_2 with $\Delta(Enc_{U_n}(m_1), Enc_{U_n}(m_2)) > 1/2$.

Proof. See blackboard.

Fact. For a random variable Y , if $E[Y] \leq \mu$ the $\Pr[Y \leq \mu] > 0$.

Let $m_1 = 0^{n+1}$, and let $S = \text{Supp}(Enc_{U_n}(m_1))$, then $|S| \leq 2^n$.

We choose a random message $m \leftarrow_R \{0, 1\}^{n+1}$ and define the following 2^n random variables for every k :

$$T_k(m) = \begin{cases} 1, & \text{if } Enc_k(m) \in S \\ 0, & \text{otherwise} \end{cases}$$

Since for every k , $Enc_k(\cdot)$ is one-to-one, we have $\Pr[T_k = 1] \leq 1/2$. Define $T = \sum_{k \in \{0,1\}^n} T_k$, then

$$E[T] = E[\sum_k T_k] = \sum_k E[T_k] \leq 2^n/2.$$

This means the probability $\Pr[T \leq 2^n/2] > 0$. In other words, there exists an m s.t. $\sum_k T_k(m) \leq 2^n/2$. For such m , **at most** half of the keys k satisfy $Enc_k(m) \in S$, i.e.,

$$\Pr[Enc_{U_n}(m) \in S] \leq 1/2.$$

Since $\Pr[Enc_{U_n}(0^{n+1}) \in S] = 1$, we have

$$\Delta(Enc_{U_n}(0^{n+1}), Enc_{U_n}(m)) > 1/2.$$



Limitation of ϵ -Statistical Security

- **Theorem 2.5** Let (Gen, Enc, Dec) be a valid encryption with $Enc : \{0, 1\}^n \times \{0, 1\}^{n+1} \rightarrow \{0, 1\}^*$. Then there exist plaintexts m_1, m_2 with $\Delta(Enc_{U_n}(m_1), Enc_{U_n}(m_2)) > 1/2$.

Proof. See blackboard.

Fact. For a random variable Y , if $E[Y] \leq \mu$ the $\Pr[Y \leq \mu] > 0$.

Let $m_1 = 0^{n+1}$, and let $S = \text{Supp}(Enc_{U_n}(m_1))$, then $|S| \leq 2^n$.

We choose a random message $m \leftarrow_R \{0, 1\}^{n+1}$ and define the following

$$T_k(m) = \begin{cases} 1, & \text{if } Enc_k(m) \in S \\ 0, & \text{otherwise} \end{cases}$$

Since for every k , $Enc_k(\cdot)$ is one-to-one, we have $\Pr[T_k = 1] \leq 1/2$. D

$$E[T] = E[\sum_k T_k] = \sum_k E[T_k] \leq 2^n/2.$$


This means the probability $\Pr[T \leq 2^n/2] > 0$. In other words, there ex

most half of the keys k satisfy $Enc_k(m) \in S$, i.e.,

$$\Pr[Enc_{U_n}(m) \in S] \leq 1/2.$$

Since $\Pr[Enc_{U_n}(0^{n+1}) \in S] = 1$, we have

$$\Delta(Enc_{U_n}(0^{n+1}), Enc_{U_n}(m)) > 1/2.$$



such m , at

Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.



Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.
 - In real life, people are using encryption with keys **shorter** than the message size to encrypt all kinds of sensitive information.
 - If the algorithm you use to break the encryption scheme runs in time 2^n , it seems **OK** since the message may be expired by then ...



Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.
 - In real life, people are using encryption with keys **shorter** than the message size to encrypt all kinds of sensitive information.
 - If the algorithm you use to break the encryption scheme runs in time 2^n , it seems **OK** since the message may be expired by then ...
- **Idea:** Would be OK if a scheme leaked information with *tiny probability* to eavesdroppers with *bounded computational resources*

Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.
 - In real life, people are using encryption with keys **shorter** than the message size to encrypt all kinds of sensitive information.
 - If the algorithm you use to break the encryption scheme runs in time 2^n , it seems **OK** since the message may be expired by then ...
- **Idea:** Would be OK if a scheme leaked information with ***tiny probability*** to eavesdroppers with ***bounded computational resources***
 - Allowing security to “**fail**” with tiny probability
 - Restricting attention to “**efficient**” attackers



Tiny probability of failure?

- Say security fails with probability 2^{-60}



Tiny probability of failure?

- Say security fails with probability 2^{-60}
 - Should we be concerned about this?

Tiny probability of failure?

- Say security fails with probability 2^{-60}
 - Should we be concerned about this?
 - With probability $> 2^{-60}$, the sender and receiver will both be struck by lightning in the next year ...
 - Something that occurs with probability $2^{-60}/\text{sec}$ is expected to occur once every **100 billion** years



Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle



Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle
 - Desktop computer $\approx 2^{57}$ keys/year
 - Supercomputer $\approx 2^{80}$ keys/year



Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle
 - Desktop computer $\approx 2^{57}$ keys/year
 - Supercomputer $\approx 2^{80}$ keys/year
 - Supercomputer since Big Bang $\approx 2^{112}$ keys
 - Restricting attention to attackers who can try 2^{112} keys is fine!



Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle
 - Desktop computer $\approx 2^{57}$ keys/year
 - Supercomputer $\approx 2^{80}$ keys/year
 - Supercomputer since Big Bang $\approx 2^{112}$ keys
 - Restricting attention to attackers who can try 2^{112} keys is fine!

Modern key space: 2^{128} keys or more ...



Computational Security

- Two problems:



Computational Security

- Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).



Computational Security

- Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

e.g., the **substitution cipher** has a huge key space, **but** can be broken efficiently.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).



Computational Security

■ Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

e.g., the **substitution cipher** has a huge key space, **but** can be broken efficiently.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).

“Breaking E is very hard”?

“Problem P cannot be solved in reasonable time”?



Computational Security

■ Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

e.g., the **substitution cipher** has a huge key space, **but** can be broken efficiently.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).

“Breaking E is very hard”?

“Problem P cannot be solved in reasonable time”?

Q: How do we **model** the resources of Eve (the adversary)?



Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.

Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.
- Let $\Pi = (Gen, Enc, Dec)$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.
- Let $\Pi = (Gen, Enc, Dec)$ be an encryption scheme with message space \mathcal{M} , and A an adversary
Define a randomized experiment $PrivK_{A, \Pi}$:
 1. A outputs $m_0, m_1 \in \mathcal{M}$
 2. $k \leftarrow Gen, b \leftarrow \{0, 1\}, c \leftarrow Enc_k(m_b)$
 3. $b' \leftarrow A(c)$

Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.
- Let $\Pi = (Gen, Enc, Dec)$ be an encryption scheme with message space \mathcal{M} , and A an adversary
Define a randomized experiment $PrivK_{A, \Pi}$:
 1. A outputs $m_0, m_1 \in \mathcal{M}$
 2. $k \leftarrow Gen, b \leftarrow \{0, 1\}, c \leftarrow Enc_k(m_b)$
 3. $b' \leftarrow A(c)$Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Perfect indistinguishability

- Let $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Define a randomized experiment $\text{PrivK}_{A,\Pi}$:

1. A outputs $m_0, m_1 \in \mathcal{M}$
2. $k \leftarrow \text{Gen}, b \leftarrow \{0, 1\}, c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.



Perfect indistinguishability

- Let $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Define a randomized experiment $\text{PrivK}_{A,\Pi}$:

1. A outputs $m_0, m_1 \in \mathcal{M}$
2. $k \leftarrow \text{Gen}, b \leftarrow \{0, 1\}, c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Π is *perfectly indistinguishable* if for *all* attackers (algorithms) A , it holds that

$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2$$

Perfect indistinguishability

- Let $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Define a randomized experiment $\text{PrivK}_{A,\Pi}$:

1. A outputs $m_0, m_1 \in \mathcal{M}$
2. $k \leftarrow \text{Gen}, b \leftarrow \{0, 1\}, c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Π is *perfectly indistinguishable* if for **all** attackers (algorithms) A , it holds that

$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2$$

Claim: Π is *perfectly indistinguishable* \Leftrightarrow Π is *perfectly secure*



Computational security?

- **Idea:** relax *perfect indistinguishability*



Computational security?

- **Idea:** relax *perfect indistinguishability*

Two approaches

- *Concrete* security
- *Asymptotic* security



Computational indistinguishability

- (t, ϵ) -indistinguishability (concrete)
 - Security may fail with probability $\leq \epsilon$
 - Restrict attention to attackers running in time $\leq t$



Computational indistinguishability

- (t, ϵ) -indistinguishability (concrete)
 - Security may fail with probability $\leq \epsilon$
 - Restrict attention to attackers running in time $\leq t$
- Π is (t, ϵ) -indistinguishable if for **all** attackers A running in time **at most** t , it holds that
$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2 + \epsilon$$



Computational indistinguishability

- (t, ϵ) -indistinguishability (concrete)
 - Security may fail with probability $\leq \epsilon$
 - Restrict attention to attackers running in time $\leq t$
- Π is (t, ϵ) -indistinguishable if for **all** attackers A running in time **at most** t , it holds that
$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2 + \epsilon$$

Does **not** lead to a **clean** theory ...

- Sensitive to exact computational model
- Π can be (t, ϵ) -secure for many choices of t, ϵ



Computational indistinguishability

- Introduce *security parameter* n (asymptotic)
 - For now, can view as the *key length*
 - Fixed by honest parties at initialization
 - Known by adversary



Computational indistinguishability

- Introduce *security parameter* n (asymptotic)
 - For now, can view as the *key length*
 - Fixed by honest parties at initialization
 - Known by adversary

Measure running time of all parties, and the *success probability* of the adversary, as *functions of n*



Computational indistinguishability

- Introduce *security parameter* n (asymptotic)
 - For now, can view as the *key length*
 - Fixed by honest parties at initialization
 - Known by adversary

Measure running time of all parties, and the *success probability* of the adversary, as *functions of n*

Computational indistinguishability:

- Security may fail with probability *negligible* in n
- Restrict attention to attackers running in time (at most) *polynomial in n*



Definitions

- A function $f : \mathbb{Z}^+ \rightarrow \mathbb{Z}^+$ is (at most) *polynomial* if **there exists** c s.t. $f(n) < n^c$ for large enough n .

A function $f : \mathbb{Z}^+ \rightarrow [0, 1]$ is *negligible* if **every** polynomial p it holds that $f(n) < 1/p(n)$ for large enough n .

- Typical example: $f(n) = \text{poly}(n) \cdot 2^{-cn}$

Why these choices?

- “Efficient” = “(probabilistic) polynomial-time (PPT)”
borrowed from *complexity theory*



Why these choices?

- “Efficient” = “(probabilistic) polynomial-time (PPT)”
borrowed from *complexity theory*
- Convenient closure properties
 - $\text{poly} * \text{poly} = \text{poly}$
 - Poly-many calls to PPT subroutine (with poly-size input) is still PPT
 - $\text{poly} * \text{negl} = \text{negl}$
 - Poly-many calls to subroutine that fails with negligible probability fails with negligible probability overall



(Re)defining encryption

- A *private-key encryption scheme* is defined by three PPT algorithms (Gen, Enc, Dec):
 - Gen: takes as input 1^n ; outputs k
 - Enc: takes as input a key k and message $m \in \{0, 1\}^*$; outputs ciphertext c : $c \leftarrow \text{Enc}_k(m)$
 - Dec: takes key k and ciphertext c as input; outputs a message m or “error” (\perp)

Computational indistinguishability (asymptotic)

■ Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}(n)$:

1. $A(1^n)$ outputs $m_0, m_1 \in \{0, 1\}^*$ of equal length
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.



Computational indistinguishability (asymptotic)

■ Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}(n)$:

1. $A(1^n)$ outputs $m_0, m_1 \in \{0, 1\}^*$ of equal length
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Π is *computationally indistinguishable* (aka *EAV-secure*) if for all PPT attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$



Example

- Consider a scheme where the **best** attack is *brute-force search* over the key space, and $Gen(1^n)$ generates a uniform n -bit key
 - So if A runs in time $t(n)$, then
$$\Pr[PrivK_{A,\Pi}(n) = 1] < 1/2 + O(t(n)/2^n)$$



Example

- Consider a scheme where the **best** attack is *brute-force search* over the key space, and $Gen(1^n)$ generates a uniform n -bit key
 - So if A runs in time $t(n)$, then
$$\Pr[PrivK_{A,\Pi}(n) = 1] < 1/2 + O(t(n)/2^n)$$
 - The scheme is **EAV-secure**: for any polynomial t , the function $t(n)/2^n$ is **negligible**.



Example

- Consider a scheme and a particular attacker A that runs for n^3 minutes and breaks the scheme with probability $2^{40}2^{-n}$
 - This does not contradict asymptotic security



Example

- Consider a scheme and a particular attacker A that runs for n^3 minutes and breaks the scheme with probability $2^{40}2^{-n}$
 - This does not contradict asymptotic security
 - What about real-world security (against this attacker)?
 - $n = 40$: A breaks with prob. 1 in 6 weeks
 - $n = 50$: A breaks with prob. 1/1000 in 3 months
 - $n = 500$: A breaks with prob. 2^{-500} in 200 years



Example

- What happens when computers get faster?
 - Consider a scheme that takes time n^2 to run but time 2^n to break with prob. 1



Example

- What happens when computers get faster?
 - Consider a scheme that takes time n^2 to run but time 2^n to break with prob. 1

What if computers get $4\times$ faster?

- Users double n ; maintain same running time
- Attacker's work is (roughly) squared!



Encryption and plaintext length

- In practice, we want encryption schemes that can encrypt **arbitrary-length** messages.



Encryption and plaintext length

- In practice, we want encryption schemes that can encrypt **arbitrary-length** messages.
- In general, encryption does **not** hide the plaintext length
 - The definition takes this into account by requiring m_0 , m_1 to have the **same** length.



Encryption and plaintext length

- In practice, we want encryption schemes that can encrypt **arbitrary-length** messages.
- In general, encryption does **not** hide the plaintext length
 - The definition takes this into account by requiring m_0 , m_1 to have the **same** length.
- But leaking plaintext length can **often** lead to problems in the real world!
 - Databases searches
 - Encrypting compressed data



Micali & Goldwasser



Silvio Micali



Shafi Goldwasser

1984: semantic security, indistinguishability
(Turing Award 2012)

Micali & Blum



Silvio Micali



Manuel Blum

1984: defined notion of pseudo-random generator
(Turing Award 1995)

Proof of Reduction




Figure 2.1: We show that the security of S' implies the security of S by transforming an adversary Eve breaking S into an adversary Eve' breaking S'

Eve breaks $S \rightarrow Eve'$ breaks S'
 S' is secure $\rightarrow S$ is secure

Next Lecture

- PRG, stream ciphers ...

