



CSE5014 CRYPTOGRAPHY AND NETWORK SECURITY

Dr. QI WANG

Department of Computer Science and Engineering

Office: Room413, CoE South Tower

Email: wangqi@sustech.edu.cn

Constructions of Authenticated Encryption

- There are three natural generic constructions:
 - Encrypt and Authenticate (**E&A**): Compute $c = Enc_k(m)$ and $t = Mac_{k_2}(m)$ and send (c, t) (SSH style)
 - Authenticate and then Encrypt (**AtE**): Compute $t = Mac_{k_2}(m)$ and then $Enc_{k_1}(t)$ (SSL style)
 - Encrypt and then Authentication (**EtA**): Compute $c = Enc_{k_1}(m)$ and $t = Mac_{k_2}(c)$ and send (c, t) (IPSec style)

Note: In all these methods, we use **independent** keys for encryption and authentication

Hash functions

- (Cryptographic) *hash function*: deterministic function mapping **arbitrary** length inputs to a short, **fixed-length** output (sometimes called a *digest*)

Hash functions

- (Cryptographic) *hash function*: deterministic function mapping **arbitrary** length inputs to a short, **fixed-length** output (sometimes called a *digest*)
- Hash functions can be *keyed* or *unkeyed*
 - In practice, hash functions are unkeyed
 - We will assume unkeyed hash functions for simplicity



Collision-resistance

- Let $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$ be a *hash function*
- A *collision* is a pair of distinct inputs x, x' such that $H(x) = H(x')$.

Collision-resistance

- Let $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$ be a *hash function*
- A *collision* is a pair of distinct inputs x, x' such that $H(x) = H(x')$.
- **Theorem 7.1** H is *collision-resistant* if it is **infeasible** to find a collision in H .

Collision-resistance

- Let $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$ be a *hash function*
- A *collision* is a pair of distinct inputs x, x' such that $H(x) = H(x')$.
- **Theorem 7.1** H is *collision-resistant* if it is **infeasible** to find a collision in H .
- What is the best “generic” collision attack on a hash function $H; \{0, 1\}^* \rightarrow \{0, 1\}^\ell$?
 - Note that collisions are **guaranteed** to exist!



Collision-resistance

- Let $H : \{0, 1\}^* \rightarrow \{0, 1\}^\ell$ be a *hash function*
- A *collision* is a pair of distinct inputs x, x' such that $H(x) = H(x')$.
- **Theorem 7.1** H is *collision-resistant* if it is **infeasible** to find a collision in H .
- What is the best “generic” collision attack on a hash function $H; \{0, 1\}^* \rightarrow \{0, 1\}^\ell$?
 - Note that collisions are **guaranteed** to exist!
 - If we compute $H(x_1), \dots, H(x_{2^\ell+1})$, we are **guaranteed** to find a collision.
 - Can we do better?



“Birthday” attacks

- Compute $H(x_1), \dots, H(x_k)$
 - What is the **probability** of a collision?

“Birthday” attacks

- Compute $H(x_1), \dots, H(x_k)$
 - What is the **probability** of a collision?
- Related to the so-called *birthday paradox*
 - How many people are needed to have a 50% chance that some two people share a birthday?



“Birthday” attacks

- Compute $H(x_1), \dots, H(x_k)$
 - What is the **probability** of a collision?
- Related to the so-called *birthday paradox*
 - How many people are needed to have a 50% chance that some two people share a birthday?

Let $n(p; H)$ be the **smallest** number of values we have to choose, such that the probability for finding a collision is **at least** p . By inverting the expression above, we have

$$n(p; H) \approx \sqrt{2H \ln \frac{1}{1-p}}.$$



“Birthday” attacks

- Compute $H(x_1), \dots, H(x_k)$
 - What is the **probability** of a collision?
- Related to the so-called *birthday paradox*
 - How many people are needed to have a 50% chance that some two people share a birthday?
- When $k \approx H^{1/2}$, probability of a collision is $\approx 50\%$
 - Birthdays: 23 people suffice!
 - Hash functions: $O(2^{\ell/2})$ hash-function evaluations



“Birthday” attacks

- Compute $H(x_1), \dots, H(x_k)$
 - What is the **probability** of a collision?
- Related to the so-called *birthday paradox*
 - How many people are needed to have a 50% chance that some two people share a birthday?
- When $k \approx H^{1/2}$, probability of a collision is $\approx 50\%$
 - Birthdays: 23 people suffice!
 - Hash functions: $O(2^{\ell/2})$ hash-function evaluations
- Need $\ell = 2n$ -bit output length to get security against attackers running in time 2^n



“Birthday bound”

- The *birthday bound* comes up in many other cryptographic contexts




“Birthday bound”

- The *birthday bound* comes up in many other cryptographic contexts
- **Example:** *IV* reuse in *CTR-mode* encryption
 - If k messages are encrypted, what is the chance that some *IV* is used **twice**?
 - Note: this is **much higher** than the probability that a *specific IV* is used again



“Birthday bound”

- The *birthday bound* comes up in many other cryptographic contexts
- **Example:** *IV* reuse in **CTR-mode** encryption



Collision resistant hash functions

■ CRH vs. PRG

- CRHs are **dual** to PRGs, in the sense that the goal is to *shrink* the input as much as possible. Similar to PRGs, if one can shrink the input by even **by one bit**, one can get a CRH collection that shrinks the bits by an amount of **polynomial**.
- In practice, people usually talk about a **single** hash func. rather than a collection. One can think of this that someone chose the key k once and then fixed the function h_k for everyone to use. In fact, most practical constructions involve some hardwired standardized constants, often known as **IV** that can be thought of as a choice of the key.



Construction of CRH

- Practical constructions of cryptographic hash functions start with a basic block, a.k.a. **compression function** $h : \{0, 1\}^{2n} \rightarrow \{0, 1\}^n$. The function $H : \{0, 1\}^* \rightarrow \{0, 1\}^n$ is defined as

$$H(m_1, \dots, m_t) = h(h(h(m_1, IV), m_2), \dots, m_t),$$

when the message is composed of t blocks (and we can pad it otherwise). This is known as the **Merkle-Damgård construction**.



Construction of CRH

- Practical constructions of cryptographic hash functions start with a basic block, a.k.a. **compression function** $h : \{0, 1\}^{2n} \rightarrow \{0, 1\}^n$. The function $H : \{0, 1\}^* \rightarrow \{0, 1\}^n$ is defined as

$$H(m_1, \dots, m_t) = h(h(h(m_1, IV), m_2), \dots, m_t),$$

when the message is composed of t blocks (and we can pad it otherwise). This is known as the **Merkle-Damgård construction**.

- **Theorem 8.1** (**Merkle-Damgård preserves collision resistance**)

Let H be constructed from h above. Then given two messages $m \neq m' \in \{0, 1\}^{tn}$ such that $H(m) = H(m')$ we can efficiently find two messages $x \neq x' \in \{0, 1\}^{2n}$ such that $h(x) = h(x')$.



Construction of CRH

■ Theorem 8.1 (Merkle-Damgård preserves collision resistance)

Let H be constructed from h above. Then given two messages $m \neq m' \in \{0, 1\}^{tn}$ such that $H(m) = H(m')$ we can efficiently find two messages $x \neq x' \in \{0, 1\}^{2n}$ such that $h(x) = h(x')$.

Proof. The intuition behind the proof is that if h is **invertible** then we could **invert** H by simply going backwards.

$$H(m_1, \dots, m_t) = h(h(h(m_1, IV), m_2), \dots, m_t),$$



Construction of CRH

■ Theorem 8.1 (Merkle-Damgård preserves collision resistance)

Let H be constructed from h above. Then given two messages $m \neq m' \in \{0, 1\}^{tn}$ such that $H(m) = H(m')$ we can efficiently find two messages $x \neq x' \in \{0, 1\}^{2n}$ such that $h(x) = h(x')$.

Proof. The intuition behind the proof is that if h is **invertible** then we could **invert** H by simply going backwards.

In principle, if a collision for H exists then so does a collision for h .

$$H(m_1, \dots, m_t) = h(h(h(m_1, IV), m_2), \dots, m_t),$$

Construction of CRH

■ Theorem 8.1 (Merkle-Damgård preserves collision resistance)

Let H be constructed from h above. Then given two messages $m \neq m' \in \{0, 1\}^{tn}$ such that $H(m) = H(m')$ we can efficiently find two messages $x \neq x' \in \{0, 1\}^{2n}$ such that $h(x) = h(x')$.

Proof. The intuition behind the proof is that if h is **invertible** then we could **invert** H by simply going backwards.

In principle, if a collision for H exists then so does a collision for h .

We look at the computation of $H(m)$ and $H(m')$ and at the first block in which the inputs differ but the output is the same (**there must be such a block**). This block will yield a collision for h .

$$H(m_1, \dots, m_t) = h(h(h(m_1, IV), m_2), \dots, m_t),$$



Hash functions in practice

■ MD5

- Developed in 1991 by Ron Rivest
- 128-bit output length
- Collisions found in 2004, and more recently in 2013, should **no longer** be used



Hash functions in practice

■ MD5

- Developed in 1991 by Ron Rivest
- 128-bit output length
- Collisions found in 2004, and more recently in 2013, should **no longer** be used

■ SHA-1

- Introduced in 1995 by NSA
- 160-bit output length
- Theoretical analysis indicates some weaknesses
- Current trend to migrate to SHA-2
- Collision found by brute force in 2017!



Hash functions in practice

■ SHA-2

- Supports 224, 256, 384, and 512-bit outputs
- No serious known weaknesses



Hash functions in practice

■ SHA-2

- Supports 224, 256, 384, and 512-bit outputs
- No serious known weaknesses

■ SHA-3 / Keccak

- Result of a public competition from 2008-2012
- Very different design from SHA-1/SHA-2
- Supports 224, 256, 384, and 512-bit outputs



Recall

- We showed how to construct a *secure MAC* for short, fixed-length messages based on any PRF/block cipher



Recall

- We showed how to construct a *secure MAC* for short, fixed-length messages based on any PRF/block cipher

Construct the following *MAC* Π :

- *Gen*: choose a uniform key k for F
- *Mac* $_k(m)$: output $F_k(m)$
- *Vrfy* $_k(m, t)$: output 1 iff $F_k(m) = t$




Recall

- We showed how to construct a *secure MAC* for short, fixed-length messages based on any PRF/block cipher
- We want to extend this to a *secure MAC* for **arbitrary-length** messages
 - Before: using *CBC-MAC*



Recall

- We showed how to construct a *secure MAC* for short, fixed-length messages based on any PRF/block cipher
- We want to extend this to a *secure MAC* for arbitrary-length messages
 - Before: using *CBC-MAC*




Recall

- We showed how to construct a *secure MAC* for short, fixed-length messages based on any PRF/block cipher
- We want to extend this to a *secure MAC* for **arbitrary-length** messages
 - Before: using *CBC-MAC*
 - After: using *hash functions*




Recall

- We showed how to construct a *secure MAC* for short, fixed-length messages based on any PRF/block cipher
- We want to extend this to a *secure MAC* for arbitrary-length messages
 - Before: using CBC-MAC
 - After: using hash functions



Hash-and-MAC



- **Theorem 8.2** If the MAC is *secure* for *fixed*-length messages and H is *collision-resistant*, then the previous construction is a *secure* MAC for *arbitrary*-length messages



- **Theorem 8.2** If the *MAC* is *secure* for *fixed*-length messages and H is *collision-resistant*, then the previous construction is a *secure* MAC for *arbitrary*-length messages
- Proof sketch
Say the sender authenticates M_1, M_2, \dots
 - Let $m_i = H(M_i)$

- **Theorem 8.2** If the *MAC* is *secure* for *fixed*-length messages and *H* is *collision-resistant*, then the previous construction is a *secure* MAC for *arbitrary*-length messages
- Proof sketch
Say the sender authenticates M_1, M_2, \dots
 - Let $m_i = H(M_i)$

Attacker outputs forgery (M, t) , $M \neq M_i$ for all i

- **Theorem 8.2** If the *MAC* is *secure* for *fixed*-length messages and *H* is *collision-resistant*, then the previous construction is a *secure* MAC for *arbitrary*-length messages

- Proof sketch

Say the sender authenticates M_1, M_2, \dots

- Let $m_i = H(M_i)$

Attacker outputs forgery (M, t) , $M \neq M_i$ for all i

Two cases:

- $H(M) = H(M_i)$ for some i
 - **Collision** in H !

- **Theorem 8.2** If the *MAC* is *secure* for *fixed*-length messages and *H* is *collision-resistant*, then the previous construction is a *secure* MAC for *arbitrary*-length messages

- Proof sketch

Say the sender authenticates M_1, M_2, \dots

- Let $m_i = H(M_i)$

Attacker outputs forgery (M, t) , $M \neq M_i$ for all i

Two cases:

- $H(M) = H(M_i)$ for some i
 - **Collision** in H !
- $H(M) \neq m_i$ for **all** i
 - **Forgery** in the underlying, *fixed*-length MAC!

Instantiation

- Hash function + block-cipher-based MAC
 - Block-length **mismatch**
 - Need to implement two crypto primitives (block cipher and hash function)



Instantiation

- Hash function + block-cipher-based MAC
 - Block-length **mismatch**
 - Need to implement two crypto primitives (block cipher and hash function)
- **HMAC**: constructed entirely from (certain type of) hash functions
 - MD5, SHA-1, SHA-2
 - **Not** SHA-3



Instantiation

- Hash function + block-cipher-based MAC
 - Block-length **mismatch**
 - Need to implement two crypto primitives (block cipher and hash function)
- **HMAC**: constructed entirely from (certain type of) hash functions
 - MD5, SHA-1, SHA-2
 - **Not** SHA-3
- Can be viewed as following the hash-and-MAC paradigm
 - With (part of the) hash function being used as a **PRF**



Other applications of hash functions

- Hash functions are ubiquitous
 - Collision-resistance \Rightarrow “fingerprinting”
 - Used as a *one-way function*
 - Used as a “random oracle”
 - Proofs of work



Other applications of hash functions

- Hash functions are ubiquitous
 - Collision-resistance \Rightarrow “fingerprinting”
 - Used as a *one-way function*
 - Used as a “random oracle”
 - Proofs of work
 - E.g., virus scanning
 - E.g., deduplication



Fingerprinting

- E.g., file integrity
 - Assuming it is possible to get a reliable copy of $H(x)$ for file x
 - Note: **different** from integrity in the context of message-authentication codes




Fingerprinting

- E.g., file integrity
 - Assuming it is possible to get a reliable copy of $H(x)$ for file x
 - Note: **different** from integrity in the context of message-authentication codes
- How to outsource files to an **untrusted** server?



x
 $h = H(x)$



$H(x) = ?h$

Outsourced storage



x_1, \dots, x_n

$h_i = H(x_i)$

$H(x_i) = ? h_i$



x_1, \dots, x_n

i

x_i

$O(n)$ client storage!

Outsourced storage



x_1, \dots, x_n

$h = H(x_1, \dots, x_n)$

$H(x_1, \dots, x_n) = ?h$



x_1, \dots, x_n

i


x_1, \dots, x_n

$O(n \cdot |x|)$ communication!

Outsourced storage




x_1, \dots, x_n




$|x_i| + O(n)$ communication!

Merkle tree



Merkle tree




$O(\log n)$ communication/computation!

Outsourced storage



x_1, \dots, x_n



$|x_i| + O(n)$ communication!

- Using a **Merkle tree**, we can solve the outsourcing problem with $O(1)$ client storage and $|x| + O(\log n)$ communication.

Key derivation

- Consider deriving a (shared) key from (shared) high-entropy information
 - E.g., biometric data
 - E.g., generating randomness



Key derivation

- Consider deriving a (shared) key from (shared) high-entropy information
 - E.g., biometric data
 - E.g., generating randomness
- Cryptographic keys must be *uniform*, but shared data is only high-entropy



Min-entropy

- Let X be a distribution
- The *min-entropy* of X (measured in bits) is
$$H_{\infty}(X) = -\log \max_x \{\Pr[X = x]\}$$
 - I.e., if $H_{\infty}(X) = n$, then the probability of guessing x sampled from X is (at most) 2^{-n}
- Min-entropy is more suitable for crypto than entropy

Min-entropy

- Let X be a distribution
- The *min-entropy* of X (measured in bits) is
$$H_{\infty}(X) = -\log \max_x \{\Pr[X = x]\}$$
 - I.e., if $H_{\infty}(X) = n$, then the probability of guessing x sampled from X is (at most) 2^{-n}
- Min-entropy is more suitable for crypto than entropy
- Given shared information x (sampled from distribution X), derive shared key $k = H(x)$
 - In what sense can we claim that k is a “good” (i.e., uniformly distributed) cryptographic key?



Private-key schemes

- We have seen how to construct schemes based on various lower-level primitives
 - Stream ciphers / PRGs
 - Block ciphers / PRFs
 - Hash functions



Private-key schemes

- We have seen how to construct schemes based on various lower-level primitives
 - Stream ciphers / PRGs
 - Block ciphers / PRFs
 - Hash functions
- How do we construct these primitives?



Two approaches

- Construct from even lower-level assumptions
 - Can prove secure (given lower-level assumption)
 - Typically **inefficient**



Two approaches

- Construct from even lower-level assumptions
 - Can prove secure (given lower-level assumption)
 - Typically **inefficient**
- Build directly
 - Much more **efficient**!
 - Need to assume security, but
 - We have formal definitions to aim for
 - We can concentrate our analysis on these primitives
 - We can develop/analyze various design principles



Terminology

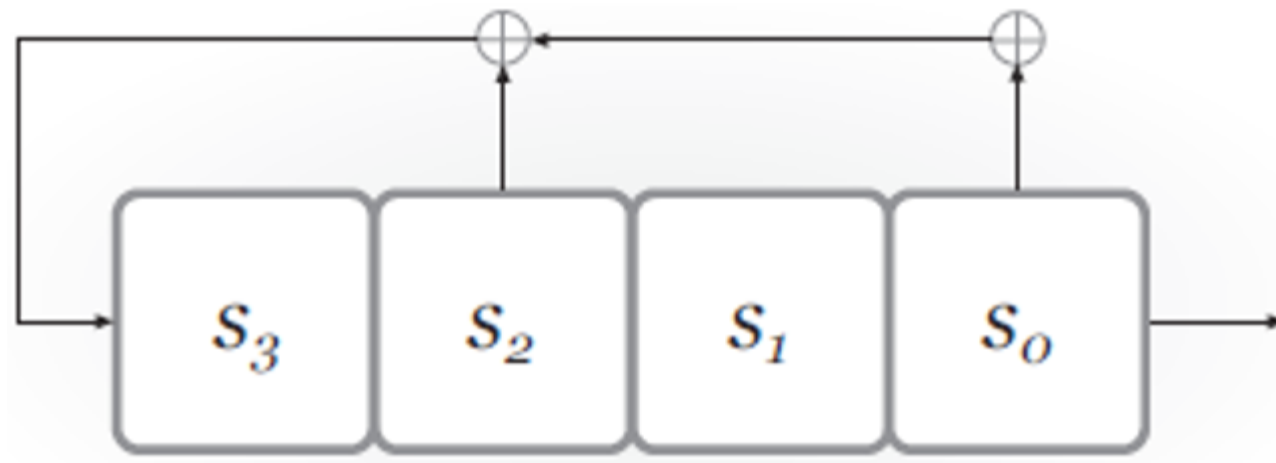
- *Init* algorithm
 - Takes as input a **key** + **initialization vector** (*IV*)
 - Outputs **initial state**
- *GetBits* algorithm
 - Takes as input the current state
 - Outputs **next** bit/byte/chunk and **updated state**
 - Allows generation of as many bits as needed

Security requirements


- If there is **no** *IV*, then (for a uniform key) the output of *GetBits* should be **indistinguishable** from a uniform, independent stream of bits
- If there is an *IV*, then (for a uniform key) the output of *GetBits* on **multiple**, uniform *IV*s should be **indistinguishable** from **multiple** uniform, independent streams of bits
 - Even if the attacker is given the *IV*s

LFSRs

- Degree $n \Rightarrow n$ registers
- State: bits s_{n-1}, \dots, s_0 (contents of the registers)
- Feedback coefficients c_{n-1}, \dots, c_0 (view as part of state; do **not** change)
- State updated and output generated in each “clock tick”



Example



- Assume initial content of registers is 0100
- First 4 state transitions:
 $0100 \rightarrow 1010 \rightarrow 0101 \rightarrow 0010 \rightarrow \dots$
- First 3 output bits:
0 0 1 \dots

LFSRs as stream ciphers

- Key + IV used to initialize the state of the LFSR (possibly including feedback coefficients)
- One bit of output per clock tick
 - State updated



LFSRs as stream ciphers

- Key + IV used to initialize the state of the LFSR (possibly including feedback coefficients)
- One bit of output per clock tick
 - State updated
- State (and output) “cycles” if state ever repeated
- *Maximal-length LFSR* cycles through all $2^n - 1$ nonzero states
 - Known how to set feedback coefficients so as to achieve maximal length



LFSRs as stream ciphers

- Key + IV used to initialize the state of the LFSR (possibly including feedback coefficients)
- One bit of output per clock tick
 - State updated
- State (and output) “cycles” if state ever repeated
- *Maximal-length LFSR* cycles through all $2^n - 1$ nonzero states
 - Known how to set feedback coefficients so as to achieve maximal length
- *Maximal-length LFSRs* have good statistical properties, **but** they are **not** cryptographically secure!

Security of LFSRs

- If feedback coefficients known ,the first n output bits directly reveal the initial state!



Security of LFSRs

- If feedback coefficients known ,the first n output bits directly reveal the initial state!
- Even if feedback coefficients are unknown, can use linear algebra to learn **everything** from $2n$ consecutive output bits (*Berlekamp-Massey algorithm*)



Security of LFSRs

- If feedback coefficients known ,the first n output bits directly reveal the initial state!
- Even if feedback coefficients are unknown, can use linear algebra to learn **everything** from $2n$ consecutive output bits (*Berlekamp-Massey algorithm*)
- Linearity is **bad** for cryptography (because linear algebra is so powerful)



Nonlinear FSRs

- Add *nonlinearity* to prevent attacks
 - Nonlinear feedback
 - Output is a nonlinear function of the state
 - Multiple (coupled) LFSRs
 - or any combination of the above




Nonlinear FSRs

- Add *nonlinearity* to prevent attacks
 - Nonlinear feedback
 - Output is a nonlinear function of the state
 - Multiple (coupled) LFSRs
 - or any combination of the above
- Still want to preserve statistical properties of the output, and **long** cycle length

Example: Trivium

- Designed by De Canniere and Preneel in 2006 as part of **eSTREAM** competition
- Intended to be simple and efficient (especially in hardware)
- Essentially **no** attacks better than **brute-force search** are known



Example: Trivium

- Three FSRs of degree 93, 84, and 111



Example: Trivium

- Three FSRs of degree 93, 84, and 111
- Initialization:
 - 80-bit key in left-most registers of first FSR
 - 80-bit IV in left-most registers of second FSR
 - Remaining registers set to 0, except for three right-most registers of third FSR
 - Run for 4×288 clock ticks



Next Lecture

- more on private-key schemes ...

