

Modern key-exchange protocols

- Security against passive eavesdroppers is **insufficient**
- Want **authenticated** key exchange
 - This requires some form of setup in advance
- Modern key-exchange protocols provide this
 - We will return to this later

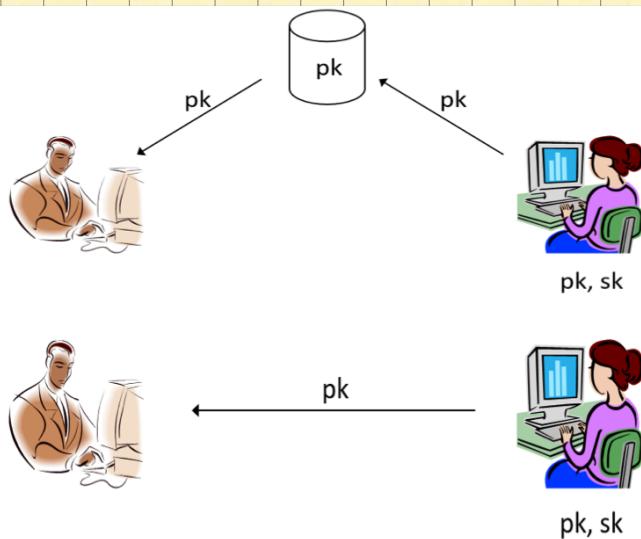
The public-key setting

- A party generates a **pair** of keys: a **public** key **pk** and a **private** key **sk**
 - Public key is widely disseminated
 - Private key is kept **secret**, and shared with no one
- Private key used by the party who generated it; public key used by everyone else
 - Also called **asymmetric cryptography**
- Security must hold even if the attacker knows **pk**

公钥公开
私钥保密

非对称加密

Public-key distribution I. II



还有 identity-based encryption:
直接用对方的 identity 作为公钥.

Public-key distribution

- Previous figures (implicitly) assume parties are able to obtain correct copies of each others' public keys
 - I.e., the attacker is **passive** during key distribution
- We will revisit this assumption later

	Private-key setting	Public-key setting
Secrecy	Private-key encryption	Public-key encryption
Integrity	Message authentication codes	Digital signature schemes

若攻击者是active的，还需要 integrity 保证密钥未被篡改。

Drawbacks of private-key crypto

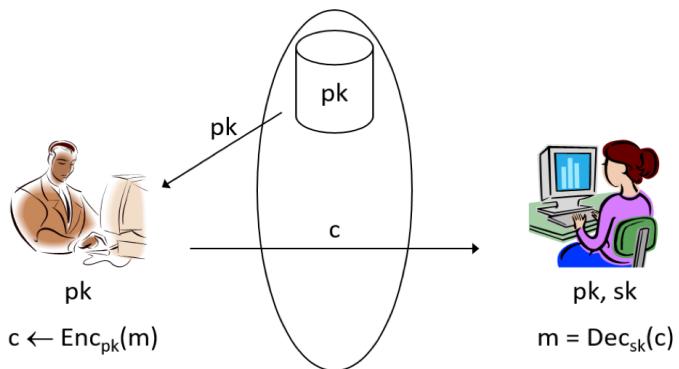
- Key distribution
 - Public keys can be distributed over *public* (but *authenticated*) channels!
- Key management in large systems of N users
 - Each user stores 1 private key and $N - 1$ public keys; only N keys overall
 - Public keys can be stored in a central directory
- Applicability in “open systems”
 - Even parties who have *no* prior relationship can find each others’ public keys and use them

Why study private-key crypto?

- Private-key cryptography is more suitable for certain applications
 - E.g., disk encryption
- Public-key crypto is roughly 2 – 3 orders of magnitude *slower* than private-key crypto
 - If private-key crypto is an option, use it!
 - Private-key crypto is used for *efficiency* even in the public-key setting

公钥加密需要更长的密钥
对称加密效率更高

Public-key encryption



- **Theorem 12.2** A *public-key encryption* scheme is composed of three PPT algorithms:

- *Gen*: *key-generation algorithm* that on input 1^n outputs pk, sk
- *Enc*: *encryption algorithm* that on input pk and a message m outputs a ciphertext c
- *Dec*: *decryption algorithm* that on input sk and a ciphertext c outputs a message m or an error ⊥

For all m and pk, sk output by *Gen*,

$$\text{Dec}_{\text{sk}}(\text{Enc}_{\text{pk}}(m)) = m$$

对公钥加密，一开始就应该设法 CPA-security (因为 pk 是公开的，相当于已经有 3 Encryption Oracle)

CPA-security

- Fix a public-key encryption scheme Π and an adversary A
- Define experiment $\text{PubK-CPA}_{A,\Pi}(n)$:
 - Run $\text{Gen}(1^n)$ to get keys pk, sk
 - Give pk to A , who outputs m_0, m_1 of same length
 - Choose uniform $b \in \{0, 1\}$ and compute the ciphertext $c \leftarrow \text{Enc}_{pk}(m_b)$; give c to A
 - A outputs a guess b' , and the experiment evaluates to 1 if $b' = b$
- Theorem 12.3** Public-key encryption scheme Π is **CPA-secure** if for all PPT adversaries A :
$$\Pr[\text{PubK-CPA}_{A,\Pi}(n) = 1] \leq 1/2 + \text{negl}(n)$$

Notes

- No encryption oracle?
 - Encryption oracle **redundant** in public-key setting
- \Rightarrow No **perfectly secret** public-key encryption
- \Rightarrow No **deterministic** public-key encryption can be CPA-secure
- \Rightarrow CPA-security implies security for encryption multiple messages as in the private-key case

公钥加密必须非完美安全以及非确定性的。

Perfectly secret public-key encryption

- Definition 12.4** A public-key encryption scheme is **perfectly secret** if for all public keys pk , all messages m_0, m_1 , all ciphertexts c , and all algorithms A , we have:
$$\Pr[A(pk, c) = 0 | c \leftarrow \text{Enc}_{pk}(m_0)] = \Pr[A(pk, c) = 0 | c \leftarrow \text{Enc}_{pk}(m_1)]$$

Theorem 12.5 No public-key encryption scheme is **perfectly secret**.

Proof.

假设 $m_0 \neq m_1$.

A : 输入 (pk, c)

for each possible choice of random bits

Set $c_0 = \text{Enc}_{pk}(m_0)$

if $c = c_0$, output 0, end

Output 1, end.

注意: ① A 肯定能停止 (有限的 random bits)

② A 可能不是 PPT 的, 但完美安全不要求 PPT.

算法正确性: ① 若 $c = \text{Enc}_{pk}(m_0)$, 则 A 输出 0

② 若 $c = \text{Enc}_{pk}(m_1)$, 则 A 输出 1

(c 不可能既是 $\text{Enc}_{pk}(m_0)$ 也是 $\text{Enc}_{pk}(m_1)$)

故 $\Pr[A(pk, c) = 0 | c \leftarrow \text{Enc}_{pk}(m_0)] = 1$ (或者反过来)

$\Pr[A(pk, c) = 0 | c \leftarrow \text{Enc}_{pk}(m_1)] = 0$

Plain RSA

- Choose random, equal-length primes p, q
- Compute modulus $N = pq$
- Choose e, d such that $e \cdot d = 1 \pmod{\phi(N)}$
- The e^{th} root of x modulo N is $x^d \pmod{N}$
 $(x^d)^e = x^{de} = x \pmod{N}$
- RSA assumption:** given N, e only, it is **hard** to compute the e^{th} root of a uniform $c \in \mathbb{Z}_N^*$

Security

- The scheme is **deterministic**
 - Cannot be CPA-secure!
- RSA assumption** only refers to hardness of computing the e^{th} roots of **uniform** c
 - c is not uniform unless m is
 - Easy to recover "small" m from c
- RSA assumption** only refers to hardness of computing the e^{th} roots **in its entirety**
 - Partial information about the e^{th} root may be leaked
- Plain RSA should **never** be used!

找出 $c^e \pmod{N}$ 是困难的
(也即 $x \pmod{N}$)

需要增加随机性

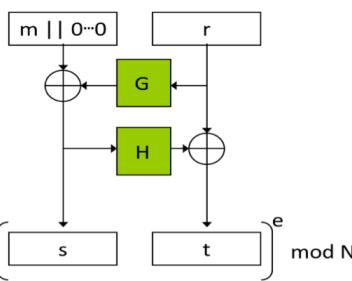
PKCS #1 v1.5

- Standard issued by RSA labs in 1993
- Idea: add **random padding**
 - To encrypt m , choose **random** r
 - $c = [(r|m)^e \pmod{N}]$
- Issues:
 - No proof of **CPA-security** (unless m is very short)
 - Chosen-plaintext attacks** known if r is too short
 - Chosen-ciphertext attacks** possible

PKCS #1 v2.0

- Optimal asymmetric encryption padding**
 - (OAEP) applied to message first
- This padding introduces **redundancy**, so that **not** every $c \in \mathbb{Z}_N^*$ is a valid ciphertext
 - Need to check for proper format upon decryption
 - Return **error** if not properly formatted

OAEPE



- RSA-OAEP can be proven *CCA-secure* under the *RSA assumption*, if G and H are modeled as *random oracles*

$$r \in \{0,1\}^{k_0}$$

$$|m| = l$$

$$|0..0| = k_1$$

$$G: \{0,1\}^{k_0} \rightarrow \{0,1\}^{l+k_1}$$

$$H: \{0,1\}^{l+k_1} \rightarrow \{0,1\}^{k_0}$$

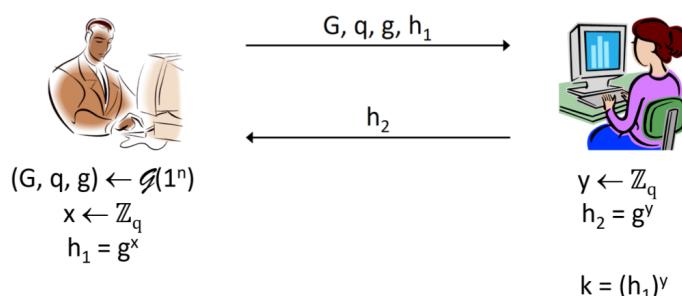
$$s = (m \parallel 0..0) \oplus G(r)$$

$$t = r \oplus H(s)$$

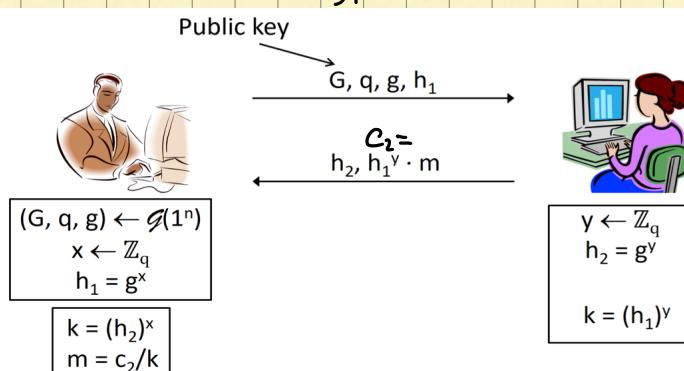
Chosen-ciphertext attacks

- Chosen-ciphertext attacks* are arguably even a greater concern in the public-key setting
 - Attacker might be a legitimate sender
 - Easier for attacker to obtain full decryptions of ciphertexts of its choice
- Related concern: *malleability*
 - I.e., given a ciphertext c that is the encryption of an unknown message m , might be possible to produce ciphertext c' that decrypts to a related message m'
 - This is also **undesirable** in the public-key setting

Diffie-Hellman key exchange



El Gamal encryption



Gen(1^n)

- Run $\mathcal{G}(1^n)$ to obtain G, q, g . Choose **uniform** $x \in \mathbb{Z}_q$. The **public key** is (G, q, g, g^x) and the **private key** is x

Enc_{pk}(m)

- where $pk = (G, q, g, h)$ and $m \in G$

Choose **uniform** $y \in \mathbb{Z}_q$. The ciphertext is $g^y, h^y \cdot m$

Dec_{sk}(c_1, c_2)

- Output c_2 / c_1^x

$$c_1 = g^y, c_2 = h^y \cdot m = g^{xy} \cdot m$$

$$c_2 / c_1^x = g^{xy} \cdot m / g^{xy} = m$$

Security

- If the *DDH assumption* is hard for \mathcal{G} , then the El Gamal encryption scheme is *CPA-secure*
 - Follows from security of Diffie-Hellman key exchange, or can be proved directly
- Dlog assumption* alone is **not** enough here

In practice

- Parameters G, q, g are standardized and shared
- Inconvenient** to treat message as group element
 - Use *key derivation* to derive a key k instead, and use k to encrypt the message
 - I.e., ciphertext is $g^y, Enc'_k(m)$, where $k = H(h^y)$
- El Gamal encryption is **not** secure against *chosen-ciphertext attacks*
 - Follows from the fact that it is *malleable*

Chosen-ciphertext attacks

- El Gamal encryption is **not** secure against *chosen-ciphertext attacks*
 - Follows from the fact that it is *malleable*
- Given ciphertext c_1, c_2 , transform it to obtain the ciphertext $c_1, c'_2 = c_1, \alpha \cdot c_2$ for **arbitrary** α
 - Since $c_1, c_2 = g^y, h^y \cdot m$, we have $c_1, c'_2 = g^y, h^y \cdot (\alpha m)$
 - I.e., encryption of m becomes an encryption of αm

Chosen-ciphertext attacks security

- Use *key derivation* coupled with *CCA-secure* private-key encryption scheme
 - I.e., ciphertext is $g^y, Enc'_k(m)$, where $k = H(h^y)$ and Enc' is a CCA-secure scheme
- Can be proved *CCA-secure* under appropriate assumptions, if H is modeled as a random oracle.

CPA-secure public key encryption

- Constructing *CCA-secure* public key encryption is **more challenging** than the private key case.
- Construction 13.1:** Construct an encryption scheme as follows:
 - Let $H : \{0,1\}^* \rightarrow \{0,1\}^n$ be a *random oracle* and $\{(f, f^{-1})\}$ be a collection of *trapdoor permutations*. The public key of the scheme will be $f(\cdot)$ while the private key is $f^{-1}(\cdot)$.
 - To *encrypt* $x \in \{0,1\}^n$, choose $r \leftarrow_R \{0,1\}^n$ and compute $f(r), H(r) \oplus x$.
 - To *decrypt* y, z , compute $r = f^{-1}(y)$ and let $x = H(r) \oplus z$.
- Theorem 13.2** The above scheme is *CPA-secure* in the random oracle model.

El Gamal 需要更强的 DDH 假设，仅有 Dlog 不够

先生成一个 key k 再加密

El Gamal 不是 CCA 安全的。

原本密文 $(c_1, c_2) = (g^y, h^y \cdot m)$

CCA 密文 $(c_1, c'_2) = (g^y, Enc'_k(m))$

$H(r) \oplus x$ 是 malleable 的，所以不是 CCA-secure

Recall

- Theorem 5.1** (*CPA security from PRFs*)

Suppose that F is a length-preserving, keyed *PRF*, then the following is a *CPA-secure encryption scheme*:

$$Enc_k(m) = \langle r, F_k(r) \oplus m \rangle$$

$$Dec_k(c_1, c_2) = c_2 \oplus F_k(c_1)$$

CPA-secure public key encryption

■ **Theorem 13.2** The above scheme is **CPA-secure** in the random oracle model.

Proof. For **public key** encryption, CPA security means that an adversary A that gets as input the encryption key $f(\cdot)$ **cannot** distinguish $Enc(x_1)$ and $Enc(x_2)$ for **every** x_1, x_2 , since encryption is public. In the random oracle model, A has access to the random oracle $H(\cdot)$.

Denote the ciphertext A gets as y^*, z^* , where $y^* = f(r^*)$ and $z^* = H(r^*) \oplus x^*$.

■ **Claim 13.2.1** The probability that A queries r^* of its oracle $H(\cdot)$ is **negl**.

Proof. Consider the following experiment: Instead of giving $z^* = H(r^*) \oplus x^*$, we give A the string $z^* = u \oplus x^*$ where u is a **uniform** element.

The only way A could tell apart the two cases is if he queries r^* to H and sees a **different** answer from u . But then we already "lost". The probability that A queries r^* in the experiment is the **same** as the probability that it queries r^* in the actual attack.

However, in this experiment, the only information A gets about r^* is $f(r^*)$. Thus, if it queries $H(\cdot)$ the value r^* , then it inverted the **trapdoor permutation**, which is almost impossible!

Proof cont'. Claim 13.2.1 means that we can ignore the probability that A queried r^* and hence we can assume that $z^* = u \oplus x^*$, where u is chosen independently at random.

However, A gets **no** information about x^* and will **not** be able to guess if it is equal to x_1 or x_2 with probability greater than $1/2$.

CCA-secure public key encryption

■ **Construction 13.3:** Construct an encryption scheme (**using two independent random oracles**) as follows:

- Let $H, H' : \{0,1\}^* \rightarrow \{0,1\}^n$ be two **independent random oracles** and $\{f, f^{-1}\}$ be a collection of **trapdoor permutations**. The public key of the scheme will be $f(\cdot)$ while the private key is $f^{-1}(\cdot)$.
- To **encrypt** $x \in \{0,1\}^n$, choose $r \leftarrow_R \{0,1\}^n$ and compute $f(r), H(r) \oplus x, H'(x, r)$.
- To **decrypt** y, z, w , compute $r = f^{-1}(y)$ and let $x = H(r) \oplus z$. Then, check that $w = H'(x, r)$: if so then return x , otherwise return \perp .

通过 $H'(x, r)$ 来验证，确保不能延展。

■ **Theorem 13.4** The above scheme is **CCA-secure** in the random oracle model.

Proof. Let A be the algorithm in a CCA attack against the scheme. Denote by y^*, z^*, w^* the challenge ciphertext A gets, where $y^* = f(r^*)$, $z^* = H(r^*) \oplus x^*$ and $w^* = H'(x^*, r^*)$.

Since H is a random oracle, we can always assume that no one (the sender, receiver, or A) can find two pairs x, r and x', r' such that $x||r \neq x'||r'$, but $H'(x, r) = H'(x', r')$.

At each step i of the attack, for every string $w \in \{0,1\}^n$, we define $H_i^{-1}(w)$ as: if the oracle H was queried before with x, r and returned w , then $H_i^{-1}(w) = (x, r)$; otherwise, $H_i^{-1}(w) = \perp$.

Observation: a pair x, r **completely determines** a ciphertext y, z, w , and y, z **completely determine** x, r .

Proof cont'. Consider the experiment: at step i , we answer a query y, z, w of A in the following way: if $H'^{-1}(w)$ is equal to some x, r that determine y, z, w , then **return x** ; otherwise, return \perp .

The **difference** between this oracle and the real decryption oracle is that we may answer \perp when the real one would give an actual answer. However, we claim that A will **not** be able to tell apart the difference with **non-negl.** probability.

The only **difference** happens if A managed to ask the oracle a query y, z, w satisfying the following:

- $w \neq w^*$.
- w was not returned as the answer of any previous query x, r to $H'(\cdot)$ by A .
- If we let x, r be the values determined by y, z , then $H'(x, r) = w$. However, since (x, r) was not asked before, the probability that 4 - this happens is only 2^{-n} .

Basically A has **no use** for the decryption box and hence it would be sufficient to prove that the scheme is just **CPA-secure**.